

# Lab 6 - Image, Datasets, Dataloaders, Augmentation

## Part I - Images

Images in Python are usually represented as ndarrays. PyTorch, of course, supports also the image as a Tensor.

Depending upon the library, images may be encoded differently:

- PIL uses, by default, the RGB encoding. The image can be represented as a  $h \times w \times 3$  ndarray by calling the `.asarray()` method of numpy
- opencv uses, by default, the BGR encoding (reverse than RGB). The image, in Python, is directly stored as a  $h \times w \times 3$  ndarray.
- PyTorch prefers images to be stored as  $3 \times h \times w$  tensors.

### Reading an image

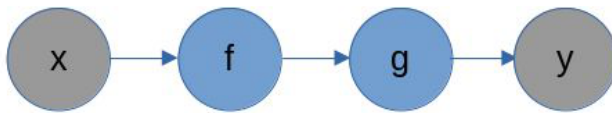
- With PIL, images can be read with the `PIL.Image.open(path)` method
- torchvision, a subpackage of PyTorch for Computer Vision, has a method `io.read_image(path)` for reading images as JPEG or PNG. For other formats, we must resort to PIL.

```
In [1]: import torchvision
import torch
from PIL import Image
import numpy as np
import os

img = Image.open("imgs/02/compgral.jpg")
print(type(img)) # the image is a type on its own
img # I can visualize the image directly like this. No need for matplotlib or other auxiliary libs
```

<class 'PIL.JpegImagePlugin.JpegImageFile'>

Out[1]:



to print the raw content of an image, we must call `np.asarray` on it.

```
In [2]: img_array = np.asarray(img)
print(img_array.shape)
img_array

(115, 472, 3)
```

```
Out[2]: array([[255, 255, 255],
               [255, 255, 255],
               [255, 255, 255],
               ...,
               [255, 255, 255],
               [255, 255, 255],
               [255, 255, 255]],

            [[255, 255, 255],
             [255, 255, 255],
             [255, 255, 255],
             ...,
             [255, 255, 255],
             [255, 255, 255],
             [255, 255, 255]],

            [[255, 255, 255],
             [255, 255, 255],
             [255, 255, 255],
             ...,
             [255, 255, 255],
             [255, 255, 255],
             [255, 255, 255]],

            ...,

            [[255, 255, 255],
             [255, 255, 255],
             [255, 255, 255],
             ...,
             [255, 255, 255],
             [255, 255, 255],
             [255, 255, 255]],

            [[255, 255, 255],
             [255, 255, 255],
             [255, 255, 255],
             ...,
             [255, 255, 255],
             [255, 255, 255],
             [255, 255, 255]],

            [[255, 255, 255],
             [255, 255, 255],
             [255, 255, 255],
             ...,
             [255, 255, 255],
             [255, 255, 255],
             [255, 255, 255]]], dtype=uint8)
```

```
In [3]: img_torch = torchvision.io.read_image("imgs/03/dataloader01.jpg")
print(type(img_torch), img_torch.shape, "\n", img_torch)
# this time the image IS a tensor
# take a look at the shape
```

```
<class 'torch.Tensor'> torch.Size([3, 319, 600])
tensor([[[255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         ...,
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255]],

        [[255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         ...,
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255]],

        [[255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         ...,
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255],
         [255, 255, 255, ..., 255, 255, 255]]], dtype=torch.uint8)
```

## Visualizing the image from a Tensor

In order to visualize the image, we can convert it to a PIL Image. But how?

```
In [4]: def tensor2PIL(tensor:torch.Tensor):

         return Image.fromarray(tensor.permute(1,2,0).numpy()) #change the order of dimensions and convert it into a tensor
         # your code here
```

## Datasets

For this lab, we will use a custom dataset for classifying cats and dogs. It is a subset of the famous dataset from the [cats vs. dogs Kaggle challenge](#). You can find it in `data/catsdogs`.

```
In [5]: folder1 = "imgs/"
        folder2 = "2"
        folder1 + "/" + folder2 # I will have a double /, so to merge use os.path.join
```

```
Out[5]: 'imgs//2'
```

```
In [6]: catsdogs_viz = [Image.open(os.path.join("data/catsdogs", im)) for im in os.listdir("data/catsdogs") if im.endswith(".
        print("Tot images", len(catsdogs_viz))
```

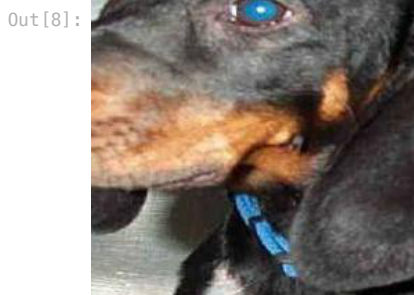
```
Tot images 10
```

Let's have a look at the data

```
In [7]: catsdogs_viz[1]
```



```
In [8]: catsdogs_viz[2]
```



## Building a custom dataset

Let us use our knowledge to build a custom dataset out of these images

```
In [9]: class CatsVsDogsDataset(torch.utils.data.Dataset):
        def __init__(self, root):
            # I don't load it as PIL, but I use torchvision.io.read_image
            self.data = [torchvision.io.read_image(os.path.join(root, im)) for im in sorted(os.listdir(root)) if im.endswith(
            #The alphabetical order of the images is not the same order in which imgs are stored in the file system!!!!
            self.labels = self._get_labels(os.path.join(root, "labels.txt"))

        def _get_labels(self, txt_path):
            with open(txt_path, "r") as f: #load the text file
                labels = [int(line.strip()) for line in f]
                #strip() is a safety method which removes whitespaces before and after something
            return labels

        def __len__(self):
            return len(self.data)

        def __getitem__(self, index):
            return self.data[index], self.labels[index]
```

now, the lazy version...

```
In [10]: class CatsVsDogsDatasetLazy(torch.utils.data.Dataset):
        def __init__(self, root):
            self.data = [os.path.join(root, im) for im in sorted(os.listdir(root)) if im.endswith(".jpg")]
            self.labels = self._get_labels(os.path.join(root, "labels.txt"))

        def _get_labels(self, txt_path):
            with open(txt_path, "r") as f:
                labels = [int(line.strip()) for line in f]
            return labels

        def __len__(self):
            return len(self.data)

        def __getitem__(self, index):
            return torchvision.io.read_image(self.data[index]), self.labels[index]
```

let us try the new dataset:

```
In [12]: dataset = CatsVsDogsDataset("data/catsdogs")

print(dataset.data)
print(dataset.labels)
print(len(dataset))
first_data = dataset[0]
print(type(first_data))
# Images are expressed by intgers from 0 to 256.
# So datatupe: dtype=torch.uint8 it's an unsigned int: must conbvert it into float!
```

```
[tensor([[[[161, 184, 196, ..., 59, 32, 21],
[139, 158, 177, ..., 44, 31, 34],
[125, 114, 119, ..., 28, 32, 46],
...,
[154, 90, 64, ..., 68, 137, 147],
[115, 167, 104, ..., 89, 162, 197],
[108, 167, 134, ..., 115, 127, 139]],

[[199, 222, 234, ..., 105, 78, 67],
[177, 196, 215, ..., 90, 77, 80],
[166, 155, 160, ..., 73, 77, 91],
...,
[177, 115, 91, ..., 111, 180, 192],
[138, 192, 131, ..., 132, 205, 242],
[131, 192, 161, ..., 156, 170, 182]],

[[160, 183, 195, ..., 68, 41, 30],
[138, 157, 176, ..., 51, 38, 41],
[126, 115, 120, ..., 34, 38, 52],
...,
[135, 75, 50, ..., 57, 126, 137],
[ 96, 152, 90, ..., 76, 149, 185],
[ 89, 152, 120, ..., 100, 114, 126]]], dtype=torch.uint8), tensor([[[[172, 173, 174, ..., 126, 134, 142],
[173, 173, 175, ..., 133, 139, 144],
[181, 182, 183, ..., 153, 155, 155],
...,
[227, 226, 223, ..., 177, 176, 176],
[228, 227, 225, ..., 177, 176, 176],
[229, 228, 225, ..., 177, 176, 176]],

[[118, 119, 119, ..., 97, 105, 114],
[119, 119, 120, ..., 102, 108, 114],
[127, 128, 128, ..., 119, 122, 122],
...,
[196, 195, 194, ..., 143, 142, 142],
[197, 196, 196, ..., 143, 142, 142],
[198, 197, 196, ..., 143, 142, 142]],

[[ 56, 57, 55, ..., 81, 87, 93],
[ 57, 57, 56, ..., 82, 87, 90],
[ 65, 66, 64, ..., 91, 91, 89],
...,
[152, 151, 152, ..., 97, 96, 96],
[153, 152, 154, ..., 97, 96, 96],
[154, 153, 154, ..., 97, 96, 96]]], dtype=torch.uint8), tensor([[[[143, 143, 143, ..., 149, 149, 149],
[143, 143, 143, ..., 148, 148, 149],
[143, 143, 143, ..., 147, 147, 147],
...,
[148, 146, 143, ..., 46, 52, 55],
[144, 142, 140, ..., 46, 52, 56],
[144, 143, 142, ..., 37, 42, 47]],

[[164, 164, 164, ..., 170, 170, 170],
[164, 164, 164, ..., 169, 169, 170],
[164, 164, 164, ..., 168, 168, 168],
...,
[116, 114, 111, ..., 44, 50, 53],
[112, 110, 108, ..., 44, 50, 54],
[112, 111, 110, ..., 35, 40, 45]],

[[193, 193, 193, ..., 199, 199, 199],
[193, 193, 193, ..., 198, 198, 199],
[193, 193, 193, ..., 197, 197, 197],
...,
[ 93, 91, 88, ..., 45, 51, 54],
[ 89, 87, 85, ..., 45, 51, 55],
[ 89, 88, 87, ..., 36, 41, 46]]], dtype=torch.uint8), tensor([[[[119, 110, 106, ..., 149, 152, 158],
[131, 126, 124, ..., 142, 145, 150],
[157, 156, 158, ..., 145, 147, 151],
...,
[130, 131, 136, ..., 126, 125, 125],
[128, 130, 132, ..., 125, 124, 123],
[129, 131, 132, ..., 123, 122, 121]],

[[128, 119, 115, ..., 152, 155, 161],
[140, 135, 133, ..., 145, 148, 153],
[166, 165, 167, ..., 148, 150, 154],
...,
[130, 131, 136, ..., 125, 124, 124],
[128, 130, 132, ..., 124, 123, 122],
[129, 131, 132, ..., 122, 121, 120]],

[[137, 128, 124, ..., 157, 160, 166],
[149, 144, 142, ..., 150, 153, 158],
[175, 174, 176, ..., 153, 155, 159],
...,
[122, 123, 128, ..., 121, 120, 120],
[120, 122, 124, ..., 120, 119, 118],
[121, 123, 124, ..., 118, 117, 116]]], dtype=torch.uint8), tensor([[[[191, 194, 193, ..., 21, 23, 24],
[187, 187, 185, ..., 19, 20, 21],
[191, 185, 182, ..., 16, 16, 17],
...,
[245, 249, 250, ..., 64, 68, 70],
```

```

[245, 248, 251, ..., 71, 72, 68],
[236, 240, 243, ..., 80, 79, 73]],

[[194, 199, 201, ..., 19, 21, 22],
[190, 192, 193, ..., 17, 18, 19],
[194, 190, 190, ..., 14, 14, 15],
...,
[251, 255, 253, ..., 64, 68, 70],
[251, 254, 254, ..., 71, 72, 68],
[242, 246, 246, ..., 80, 79, 73]],

[[201, 203, 203, ..., 20, 22, 23],
[197, 196, 195, ..., 18, 19, 20],
[201, 194, 192, ..., 15, 15, 16],
...,
[241, 245, 244, ..., 64, 68, 70],
[241, 244, 245, ..., 71, 72, 68],
[232, 236, 237, ..., 80, 79, 73]], dtype=torch.uint8), tensor([[[120, 120, 120, ..., 122, 121, 121],
[120, 120, 120, ..., 122, 121, 121],
[120, 120, 120, ..., 122, 121, 121],
...,
[113, 113, 113, ..., 102, 96, 91],
[113, 113, 113, ..., 105, 99, 91],
[113, 113, 113, ..., 107, 101, 91]],

[[114, 114, 114, ..., 113, 112, 112],
[114, 114, 114, ..., 113, 112, 112],
[114, 114, 114, ..., 113, 112, 112],
...,
[109, 109, 109, ..., 102, 96, 92],
[109, 109, 109, ..., 105, 99, 92],
[109, 109, 109, ..., 107, 101, 92]],

[[ 56, 54, 54, ..., 74, 73, 73],
[ 56, 54, 54, ..., 74, 73, 73],
[ 56, 54, 52, ..., 74, 73, 73],
...,
[ 72, 72, 72, ..., 90, 84, 78],
[ 72, 72, 72, ..., 93, 87, 78],
[ 72, 72, 72, ..., 95, 89, 78]], dtype=torch.uint8), tensor([[[209, 137, 107, ..., 190, 111, 21],
[205, 146, 108, ..., 148, 137, 41],
[210, 192, 163, ..., 64, 108, 74],
...,
[160, 169, 187, ..., 165, 159, 142],
[156, 169, 191, ..., 174, 147, 131],
[168, 184, 203, ..., 182, 138, 124]],

[[211, 139, 109, ..., 188, 111, 21],
[206, 148, 110, ..., 146, 137, 41],
[211, 193, 164, ..., 62, 108, 74],
...,
[159, 168, 186, ..., 166, 160, 143],
[155, 168, 187, ..., 175, 148, 132],
[164, 180, 199, ..., 183, 139, 125]],

[[200, 128, 98, ..., 176, 99, 9],
[198, 137, 99, ..., 134, 125, 29],
[203, 185, 156, ..., 50, 96, 62],
...,
[155, 164, 182, ..., 161, 155, 138],
[151, 164, 184, ..., 170, 143, 127],
[161, 177, 196, ..., 178, 134, 120]], dtype=torch.uint8), tensor([[[182, 181, 180, ..., 196, 198, 198],
[182, 181, 180, ..., 196, 197, 197],
[182, 182, 181, ..., 196, 196, 196],
...,
[230, 231, 234, ..., 182, 194, 198],
[231, 232, 234, ..., 173, 188, 195],
[230, 231, 233, ..., 148, 164, 173]],

[[183, 182, 181, ..., 193, 195, 195],
[183, 182, 181, ..., 193, 194, 194],
[183, 183, 182, ..., 193, 193, 193],
...,
[231, 232, 235, ..., 183, 195, 199],
[232, 233, 235, ..., 174, 189, 196],
[231, 232, 234, ..., 149, 165, 174]],

[[175, 174, 173, ..., 174, 178, 178],
[175, 174, 173, ..., 176, 177, 179],
[175, 175, 174, ..., 176, 176, 178],
...,
[223, 224, 227, ..., 175, 187, 191],
[224, 225, 227, ..., 166, 181, 188],
[223, 224, 226, ..., 141, 157, 166]], dtype=torch.uint8), tensor([[[217, 217, 217, ..., 52, 54, 56],
[217, 217, 217, ..., 55, 54, 53],
[217, 217, 217, ..., 56, 52, 48],
...,
[216, 215, 215, ..., 9, 9, 9],
[214, 214, 215, ..., 9, 9, 9],
[214, 214, 215, ..., 9, 9, 9]],

[[204, 204, 204, ..., 47, 49, 51],
[204, 204, 204, ..., 50, 49, 48],

```

```

[204, 204, 204, ..., 51, 47, 43],
...,
[207, 206, 206, ..., 9, 9, 9],
[205, 205, 206, ..., 9, 9, 9],
[205, 205, 206, ..., 9, 9, 9]],
dtype=torch.uint8), tensor([[[ 74, 71, 67, ..., 73, 75, 78],
[ 71, 68, 67, ..., 70, 72, 75],
[ 71, 71, 72, ..., 73, 74, 76],
...,
[254, 249, 249, ..., 51, 61, 45],
[252, 247, 251, ..., 52, 68, 57],
[252, 252, 255, ..., 51, 77, 71]],
dtype=torch.uint8))
[0, 1, 0, 0, 0, 1, 1, 1, 1, 0]
10
<class 'tuple'>

```

## Part II - DataLoaders

DataLoaders can be quickly constructed from a Dataset...

```
In [13]: dataloader = torch.utils.data.DataLoader(dataset, batch_size=2, shuffle=True) #non lazy version
```

let us loop through the dataloader:

```
In [14]: list_of_images = []
for imgs, labels in dataloader:
    # Dataloader is an iterator: yield operator!
    # Each query returns a different item in the dataloader, and then it stitches them together. (collate function)
    print(imgs.shape, labels)
    for img in imgs:
        list_of_images.append(tensor2PIL(img))

torch.Size([2, 3, 216, 237]) tensor([0, 0])
torch.Size([2, 3, 216, 237]) tensor([1, 1])
torch.Size([2, 3, 216, 237]) tensor([1, 0])
torch.Size([2, 3, 216, 237]) tensor([0, 0])
torch.Size([2, 3, 216, 237]) tensor([1, 1])
```

**Q:** Do you notice something different w.r.t. what we saw before?

```
In [15]: list_of_images[0] # I have shuffled them so I have a different order
```

Out[15]:



```
In [16]: list_of_images[4]
```

Out[16]:



## Part III - Data Augmentation

From the lecture, we know that we can construct custom augmentation pipelines using torchvision.

Let us quickly implement augmentability in our (non-lazy) dataset

```
In [ ]: class CatsVsDogsDatasetAugmentable(CatsVsDogsDataset): # the len method is inherited by CatsVsDogsDataset
    def __init__(self, root, transform=None):
        super().__init__(root)
        self.transform = transform

    def __getitem__(self, index):
        img, label = super().__getitem__(index)
        if self.transform:
            img = self.transform(img)
        return img, label
```

Q: What about the `__len__` method?

```
In [ ]: transform_pipeline = torchvision.transforms.Compose([
    torchvision.transforms.RandomHorizontalFlip(),
    torchvision.transforms.RandomRotation(degrees=15),
])

aug_dataset = CatsVsDogsDatasetAugmentable("data/catsdogs", transform=transform_pipeline)
```

```
In [ ]: tensor2PIL(aug_dataset.data[0])
```

```
In [ ]: tensor2PIL(aug_dataset[0][0])
```

Q: What is the difference between the two code snippets here above?

Problem to solve:

- Suppose we have a dataset composed of  $n$  images
- As opposed to the cats vs dogs example we saw before, the images don't have a common size  $h \times w$
- **What are my possibilities for training an ANN on this dataset?**
  - crop imgs
  - batch size of 1 but and then effective batch size bigger, so each time I accumulate the gradient and after reaching the desired batch size.

## End of "compulsory" lab. Next we have some optional suggestions for loading datasets doing and data augmentation in PyTorch

The "compulsory" augmentations: how to do them on our dataset

```
pipeline = T.Compose([
    ...,
    T.ToTensor(),
    T.Normalize(mean=mean, std=std)
])
```

Essentially, we need to calculate `mean` and `std`.

We have our data in `dataset.data`. What can we do to get mean and std?

```
In [ ]: # your code here
```

## Train/Test splitting in PyTorch

We can apply a train/test split by using the `torch.utils.data.random_split` method

```
In [ ]: pct_train = .7
len_train = int(len(aug_dataset) * pct_train)
```



```
len_test = len(aug_dataset) - len_train
# don't do:
# len_test = int(len(aug_dataset) * (1-pct_train))
# for casting reasons: casting to integers = casting to floor!
trainset, testset = torch.utils.data.random_split(aug_dataset, [len_train, len_test])
```

Notice that now, the `trainset` and `aug_dataset` are of two different types!

```
In [ ]: print(type(aug_dataset), type(trainset), type(testset))
```

we can recover the original dataset by accessing the `dataset` attribute of `torch.utils.data.dataset.Subset`

```
In [ ]: trainset.dataset.labels
```

Despite the difference, both `trainset` and `aug_dataset` can be equally used to create DataLoaders...

## Miscellaneous dataset helps

### ImageFolder

It often happens that datasets are distributed with the following folder structure:

```
root_folder
|
- class 0
  |
  - images belonging to class 0
  |
  ...
  |
- class i
  |
  - images belonging to class i
  |
  ...
```

without a corresponding `labels.txt` file (or similar file.)

When the situation is this one, without building exotic custom classes, we can use the `torchvision.datasets.ImageFolder(...)` class that automatically builds a (lazy) dataset for us.

## Downloading widespread benchmark datasets

To download benchmark datasets like

- MNIST
- Cifar10 and Cifar100
- Fashion-MNIST
- Microsoft COCO
- ...

we can use the corresponding `torchvision.datasets` classes. Just a couple of notes:

- remember to set, preferably, the flag `download` to True in the constructor (otherwise the dataset won't download)
- ImageNet won't download because of recent controversies on fairness and privacy. If you need it, download it (at your home) from [here](#)

## References

- Pillow docs: <https://pillow.readthedocs.io/en/stable/>
- torch tutorial on datasets and dataloaders: [https://pytorch.org/tutorials/beginner/basics/data\\_tutorial.html#datasets-dataloaders](https://pytorch.org/tutorials/beginner/basics/data_tutorial.html#datasets-dataloaders)
- torchvision tutorials and docs
  - IO: <https://pytorch.org/vision/0.8/io.html>
  - Datasets: <https://pytorch.org/vision/stable/datasets.html>
  - Transformations: <https://pytorch.org/vision/stable/transforms.html>

### Additional material

- Albumentations, library for more advanced data augmentation: <https://albumentations.ai/>